



# Algorithms and Applications of Data Mining

Yijun Lin  
[yijunlin@usc.edu](mailto:yijunlin@usc.edu)

01/31



# About This Course

- Spring 2021, Friday, 6-8 PM PST

- Instructor: Yao-Yi Chiang

- TA: Yijun Lin
  - Office Hour Sat. 7-9 PM PST

- Syllabus:

	Topic	Readings and Assignments	Deliverables/Due Dates
<b>Week 1</b>	Introduction to Data Mining	Ch1: Data Mining and	
<b>Week 2</b>	MapReduce	Ch2: Large-Scale File Systems and Map-Reduce	Homework 1 assigned
<b>Week 3</b>	Frequent itemsets and Association rules	Ch6: Frequent itemsets,	Homework 2 assigned
<b>Week 4</b>	Clustering	Ch7: Clustering	Homework 1 due
<b>Week 5</b>	Recommendation Systems: Content-based	Ch9: Recommendation systems	Homework 2 due, Homework 3 assigned
<b>Week 6</b>	Recommendation Systems: Collaborative Filtering	Ch9: Recommendation systems	Homework 3 due



# Assignments

- Theoretical and programming questions
  - Real-world datasets
- Homework 1 - basic spark operations
- Homework 2 - mining frequent itemset
- Homework 3 - recommender system
- Optional – clustering



# Config Environment

- **Python** is required for all the assignments
- Implementing with Apache Spark Framework
  - python=3.7
  - pyspark=3.0.1
  - git clone <https://github.com/linyijun/cis-data-mining-ta-materials>
- Install miniconda/anaconda
  - conda env create -f spark-env.yml python=3.7
- Install PyCharm



# Introduction to Spark

Thanks for source slide and material to: Dr. Heather Miller  
<https://www.coursera.org/learn/scala-spark-big-data/home/welcome>

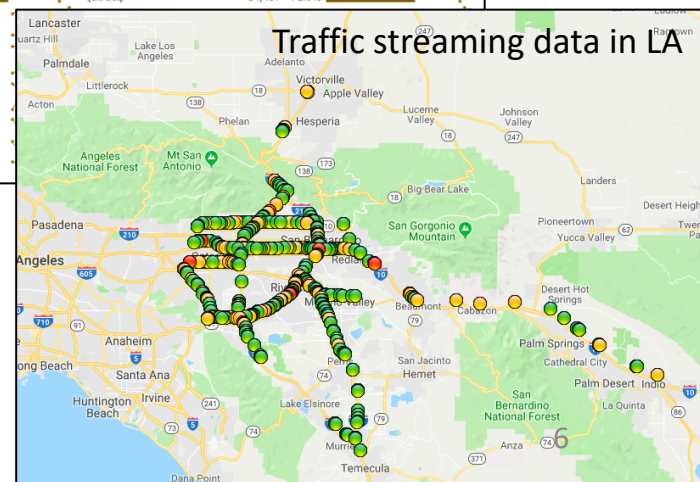
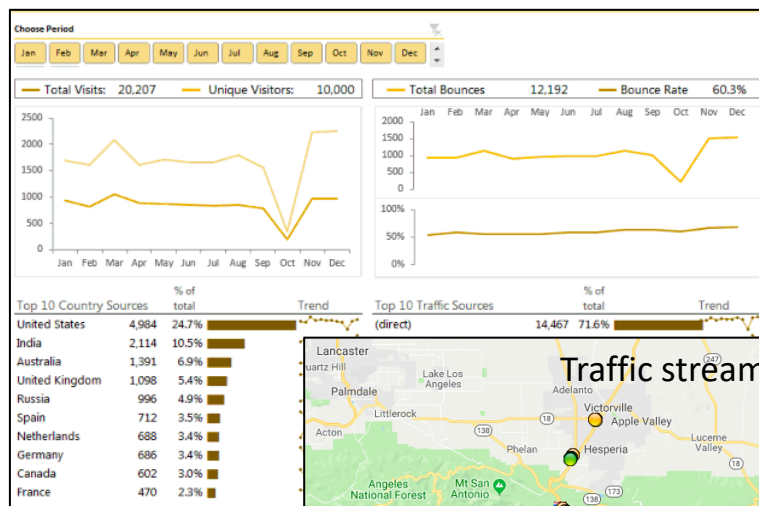


# What is Spark?



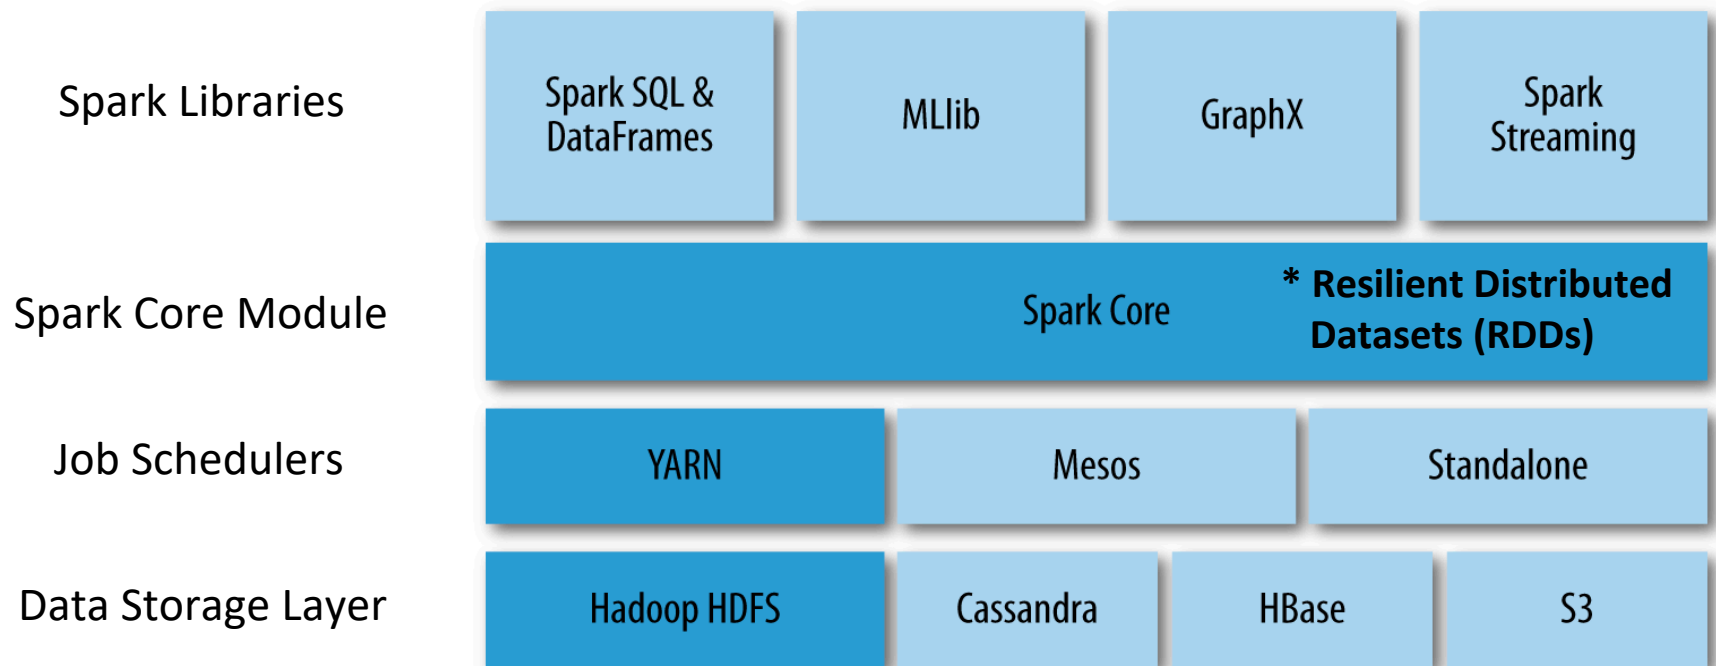
- **Apache Spark** is a unified analytics engine for large-scale data processing

- Application areas
  - Interactive Data Query
  - Real-time Data Analysis
  - Streaming Data Processing






# Spark Stack





# Resilient Distributed Datasets (RDDs)

- An RDD is an **immutable, in-memory collection** of objects
- Each RDD can be split into multiple partitions, which in turn are computed on different nodes of the cluster
- RDDs seem a lot like Scala collections
  - `RDD[T]` and `List[T]`







# How to create an RDD

- RDDs can be created in two ways:
  - Creating from a SparkContext object
  - Transforming from an existing RDD

# How to create an RDD (Cont.)



- Creating from a **SparkContext** object
  - Can be thought as your handle to the Spark cluster
  - Represents the connection to a Spark cluster

```
if __name__ == '__main__':  
  
    sc_conf = pyspark.SparkConf() \  
        .setAppName('task1') \  
        .setMaster('local[*]') \  
        .set('spark.driver.memory', '8g') \  
        .set('spark.executor.memory', '4g')  
  
    sc = pyspark.SparkContext(conf=sc_conf)  
    sc.setLogLevel("OFF")
```

# How to create an RDD (Cont.)



- Creating from a **SparkContext** object
  - **parallelize**: convert a local Scala collection to an RDD

```
a_list = ['you', 'jump', 'I', 'jump', '']  
a_rdd = sc.parallelize(a_list) # RDD[String]
```

# How to create an RDD (Cont.)



- Creating from a **SparkContext** object
  - **parallelize**: convert a local Scala collection to an RDD

```
a_list = ['you', 'jump', 'I', 'jump', '']  
a_rdd = sc.parallelize(a_list) # RDD[String]
```

- **textFile**: read a file from HDFS or local file system

```
input_file = 'work-count-sample-doc.txt'  
text_rdd = sc.textFile(input_file)
```

# How to create an RDD (Cont.)



- Transforming from an existing RDD
  - E.g., calling a *map operation* on an existing RDD, it will return a new RDD

```
# call a map operation on an RDD
length_rdd = word_rdd.map(lambda x: len(x)) # RDD[Int]
```



# RDD Operations

- Transformations
  - E.g., map, filter, ...

```
# call a map operation on an RDD
length_rdd = word_rdd.map(lambda x: len(x)) # RDD[Int]
```

- Actions
  - E.g., collect, reduce ...

```
a_coll = a_rdd.collect() # RDD -> collection
print(a_coll) # ['you', 'jump', 'I', 'jump', '']
```



# Transformations VS Actions

- Transformations
  - Return new RDDs as results
  - They are **lazy**, the result RDD is **not immediately computed**
- Actions
  - Compute a result based on an RDD, and returned
  - They are **eager**, the result is **immediately computed**



# Transformations VS Actions

- Transformations

- Return new RDDs as results
- They are **lazy**, the result RDD is **not immediately computed**

```
# call a map operation on an RDD
length_rdd = word_rdd.map(lambda x: len(x)) # RDD[Int]
```

- Actions

- Compute a result based on an RDD, and returned
- They are **eager**, the result is **immediately computed**

```
a_coll = a_rdd.collect() # RDD -> collection
print(a_coll) # ['you', 'jump', 'I', 'jump', '']
```



# Word Count

```
word_count.py x text.txt x
1 from pyspark import SparkContext
2 import os
3
4 os.environ['PYSPARK_PYTHON'] = '/usr/local/bin/python3.6'
5 os.environ['PYSPARK_DRIVER_PYTHON'] = '/usr/local/bin/python3.6'
6
7 sc = SparkContext('local[*]', 'wordCount')
8
9 input_file_path = './text.txt'
10 textRDD = sc.textFile(input_file_path)
11
12 counts = textRDD.flatMap(lambda line: line.split(' ')) \
13     .map(lambda word: (word, 1)).reduceByKey(lambda a, b: a+b).collect()
14
15 for each_word in counts:
16     print(each_word)
17
```



# If you want to learn more...

- Official documentation
  - <http://spark.apache.org/docs/latest/>
- Online course
  - Coursera: Big Data Analysis with Scala and Spark
- Books
  - *Learning Spark, O' Reilly*
  - *Advanced Analytics with Spark: Patterns for Learning from Data at Scale, O' Reilly*
  - *Machine Learning with Spark, Packt*