



Algorithms and Applications of Data Mining

Yijun Lin

yijunlin@usc.edu

02/06

Map-Reduce

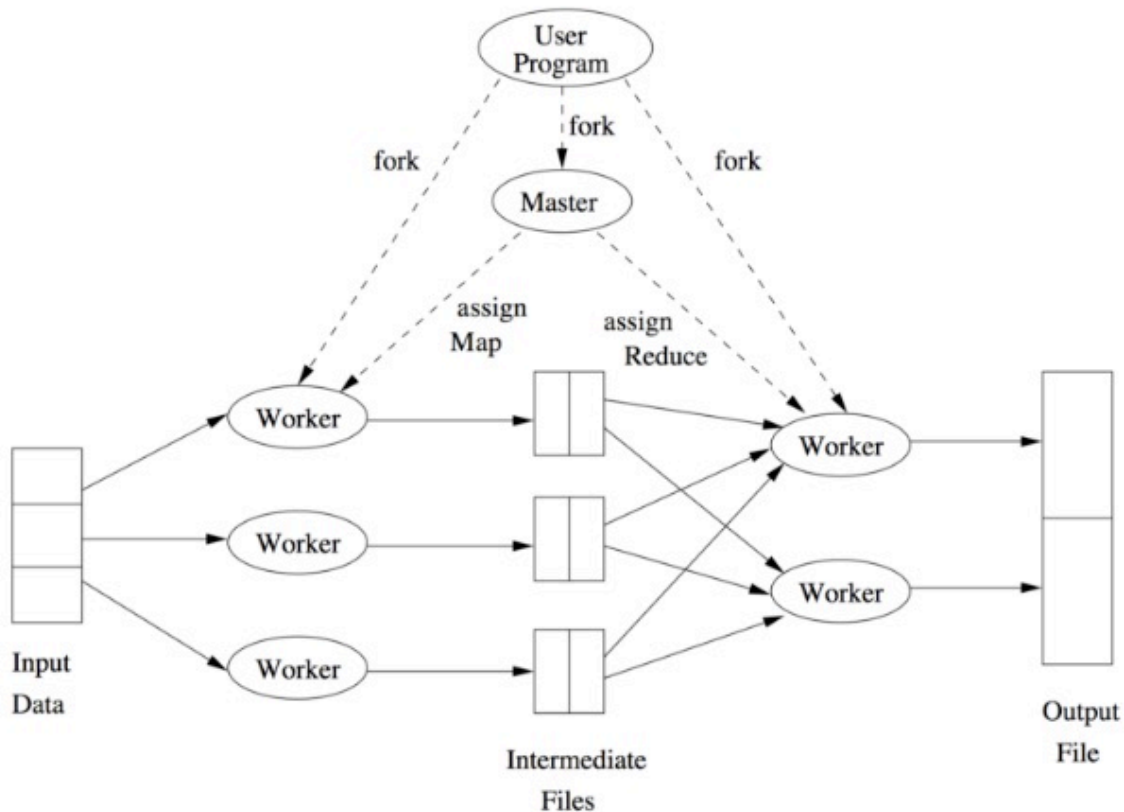



Figure 2.3: Overview of the execution of a MapReduce program

Q1. Concerning the above figure, where does the MapReduce program store 1. Input data, 2. Intermediate files, and 3. Output data?

Map-Reduce



1. DFS/HDFS
2. Local FS
3. DFS/HDFS

Figure 2.3: Overview of the execution of a MapReduce program

Q1. Concerning the above figure, where does the MapReduce program store 1. Input data, 2. Intermediate files, and 3. Output data?

Map-Reduce

Write the Map and Reduce tasks and their output for joining these two tables:

Order(orderid, account, date)

1, aaa, d1
2, aaa, d2
3, bbb, d3

LineItem(orderid, itemid, qty)

1, 10, 1
1, 20, 3
2, 10, 5
2, 50, 100
3, 20, 1

Map-Reduce

Map task:

Use orderid as key, and table name together with other columns as value. Map each row in the table and emit the key-value pair.

	relation name
Order	
1, aaa, d1	→ 1 : "Order", (1,aaa,d1)
2, aaa, d2	→ 2 : "Order", (2,aaa,d2)
3, bbb, d3	→ 3 : "Order", (3,bbb,d3)
Line	
1, 10, 1	→ 1 : "Line", (1, 10, 1)
1, 20, 3	→ 1 : "Line", (1, 20, 3)
2, 10, 5	→ 2 : "Line", (2, 10, 5)
2, 50, 100	→ 2 : "Line", (2, 50, 100)
3, 20, 1	→ 3 : "Line", (3, 20, 1)

Reduce task:

groups together all values (tuples) associated with each key and emit joined values.

Reducer for key 1

"Order", (1,aaa,d1)
"Line", (1, 10, 1)
"Line", (1, 20, 3)



(1, aaa, d1, 1, 10, 1)
(1, aaa, d1, 1, 20, 3)




Introduction to Spark

Thanks for source slide and material to: Dr. Heather Miller
<https://www.coursera.org/learn/scala-spark-big-data/home/welcome>



Resilient Distributed Datasets (RDDs)

- An RDD is an **immutable, in-memory collection** of objects
- Each RDD can be split into multiple partitions, which in turn are computed on different nodes of the cluster
- RDDs seem a lot like Scala collections
 - `RDD[T]` and `List[T]`





Transformations VS Actions

- Transformations
 - Return new RDDs as results
 - They are **lazy**, the result RDD is **not immediately computed**
- Actions
 - Compute a result based on an RDD, and returned
 - They are **eager**, the result is **immediately computed**



Transformations VS Actions

- Transformations

- Return new RDDs as results
- They are **lazy**, the result RDD is **not immediately computed**

```
# call a map operation on an RDD
length_rdd = word_rdd.map(lambda x: len(x)) # RDD[Int]
```

- Actions

- Compute a result based on an RDD, and returned
- They are **eager**, the result is **immediately computed**

```
a_coll = a_rdd.collect() # RDD -> collection
print(a_coll) # ['you', 'jump', 'I', 'jump', '']
```



Transformations VS Actions

- Transformations
 - Return new RDDs as results
 - They are **lazy**, the result RDD is **not immediately computed**
- Actions
 - Compute a result based on an RDD, and returned
 - They are **eager**, the result is **immediately computed**

Laziness / Eagerness is how we can limit network communication using the programming model



Example

- Consider the following example:

```
a_list = ['you', 'jump', 'I', 'jump', '']  
# create an RDD from a list  
a_rdd = sc.parallelize(a_list) # RDD[String]  
# call a map operation RDD  
a_len_rdd = a_rdd.map(lambda x: len(x)) # RDD[Int]
```

What has happened on the cluster at this point?





Example (Cont.)

- Consider the following example:

```
a_list = ['you', 'jump', 'I', 'jump', '']  
# create an RDD from a list  
a_rdd = sc.parallelize(a_list) # RDD[String]  
# call a map operation RDD  
a_len_rdd = a_rdd.map(lambda x: len(x)) # RDD[Int]
```

What has happened on the cluster at this point?

Nothing. Execution of map (a transformation) is deferred.



Example (Cont.)

- Consider the following example:

```
a_list = ['you', 'jump', 'I', 'jump', '']  
# create an RDD from a list  
a_rdd = sc.parallelize(a_list) # RDD[String]  
# call a map operation RDD  
a_len_rdd = a_rdd.map(lambda x: len(x)) # RDD[Int]
```

What has happened on the cluster at this point?

Nothing. Execution of *map* (a transformation) is deferred.

How to ensure this computation is done on the cluster?



Example (Cont.)

- Consider the following example:

```
a_list = ['you', 'jump', 'I', 'jump', '']  
# create an RDD from a list  
a_rdd = sc.parallelize(a_list) # RDD[String]  
# call a map operation RDD  
a_len_rdd = a_rdd.map(lambda x: len(x)) # RDD[Int]  
  
total_len = a_len_rdd.reduce(lambda a, b: a + b) # 12
```

add an action, *reduce*

Spark starts the execution when an action is called

Return the total number of characters in the entire RDD of strings



Benefits of Laziness

- Another example:

```
input_file = 'work-count-sample-doc.txt'
text_rdd = sc.textFile(input_file)
word_rdd = text_rdd.flatMap(lambda x: x.split(' ')).take(10)
```



Benefits of Laziness

- Another example:

```
input_file = 'work-count-sample-doc.txt'
text_rdd = sc.textFile(input_file)
word_rdd = text_rdd.flatMap(lambda x: x.split(' ')).take(10)
```

- The execution of *filter* is **deferred** until the *take* action happens
 - Spark will not compute intermediate RDDs. As soon as 10 elements of the filtered RDD have been computed, *logsWithErrors* is done.



Benefits of Laziness

- Another example:

```
input_file = 'work-count-sample-doc.txt'
text_rdd = sc.textFile(input_file)
word_rdd = text_rdd.flatMap(lambda x: x.split(' ')).take(10)
```

- The execution of *filter* is **deferred** until the *take* action happens
 - Spark will not compute intermediate RDDs. As soon as 10 elements of the filtered RDD have been computed, *logsWithErrors* is done
- Spark leverages this by analyzing and optimizing the **chain of operations** before executing it
 - Spark saves time and space to compute elements of the unused result of the *filter operation*



Common Transformations

map **map[T](f: A=>B): RDD[T]**

Apply function to each element in the RDD and return an RDD of the result.

flatMap **flatMap[T](f: A=>B): RDD[T]**

Apply function to each element in the RDD and return an RDD of the result, but output is flattened.

filter **filter[T](pred: A=>Boolean): RDD[T]**

Apply predicate function, pred, to each element in the RDD and return an RDD of elements that passed the condition.

distinct **distinct():RDD[T]**

Return an RDD with duplicates removed



Common Transformations

flatMap **flatMap[T](f: A=>B): RDD[T]**
Apply function to each element in the RDD and return an RDD of the result, but output is flattened.



Common Transformations

distinct **distinct():RDD[T]**
Return an RDD with duplicates removed



Common Actions

- collect** **collect: Array[T]**
Return all elements from RDD.
- count** **count(): Long**
Return the number of elements in the RDD.
- take** **take(num: Int): Array[T]**
Return the first num elements of the RDD.
- reduce** **reduce(op: (A, A) => A): A**
Combine the elements in the RDD together using op function and return result.
- foreach** **foreach(f: A => Unit): Unit**
Apply function to each element in the RDD, and return Unit.



Common Actions

count **count(): Long**
Return the number of elements in the RDD.



Common Actions

foreach **foreach(f: A => Unit): Unit**

Apply function to each element in the RDD, and return Unit.

```
case class Person(name: String, age: Int)

val people: RDD[Person]
people.foreach(println)
```



Cluster Topology

- A simple example with *println*

```
case class Person(name: String, age: Int)

val people: RDD[Person]
people.foreach(println)
```


What happens?

How Spark jobs are Executed




Master-Worker
Topology

Master



Workers



How Spark jobs are Executed






This is the node you're interacting with when you're writing Spark programs!



In the context of a Spark program


These are the nodes actually executing the jobs!



How Spark jobs are Executed




This is the node you're interacting with when you're writing Spark programs!



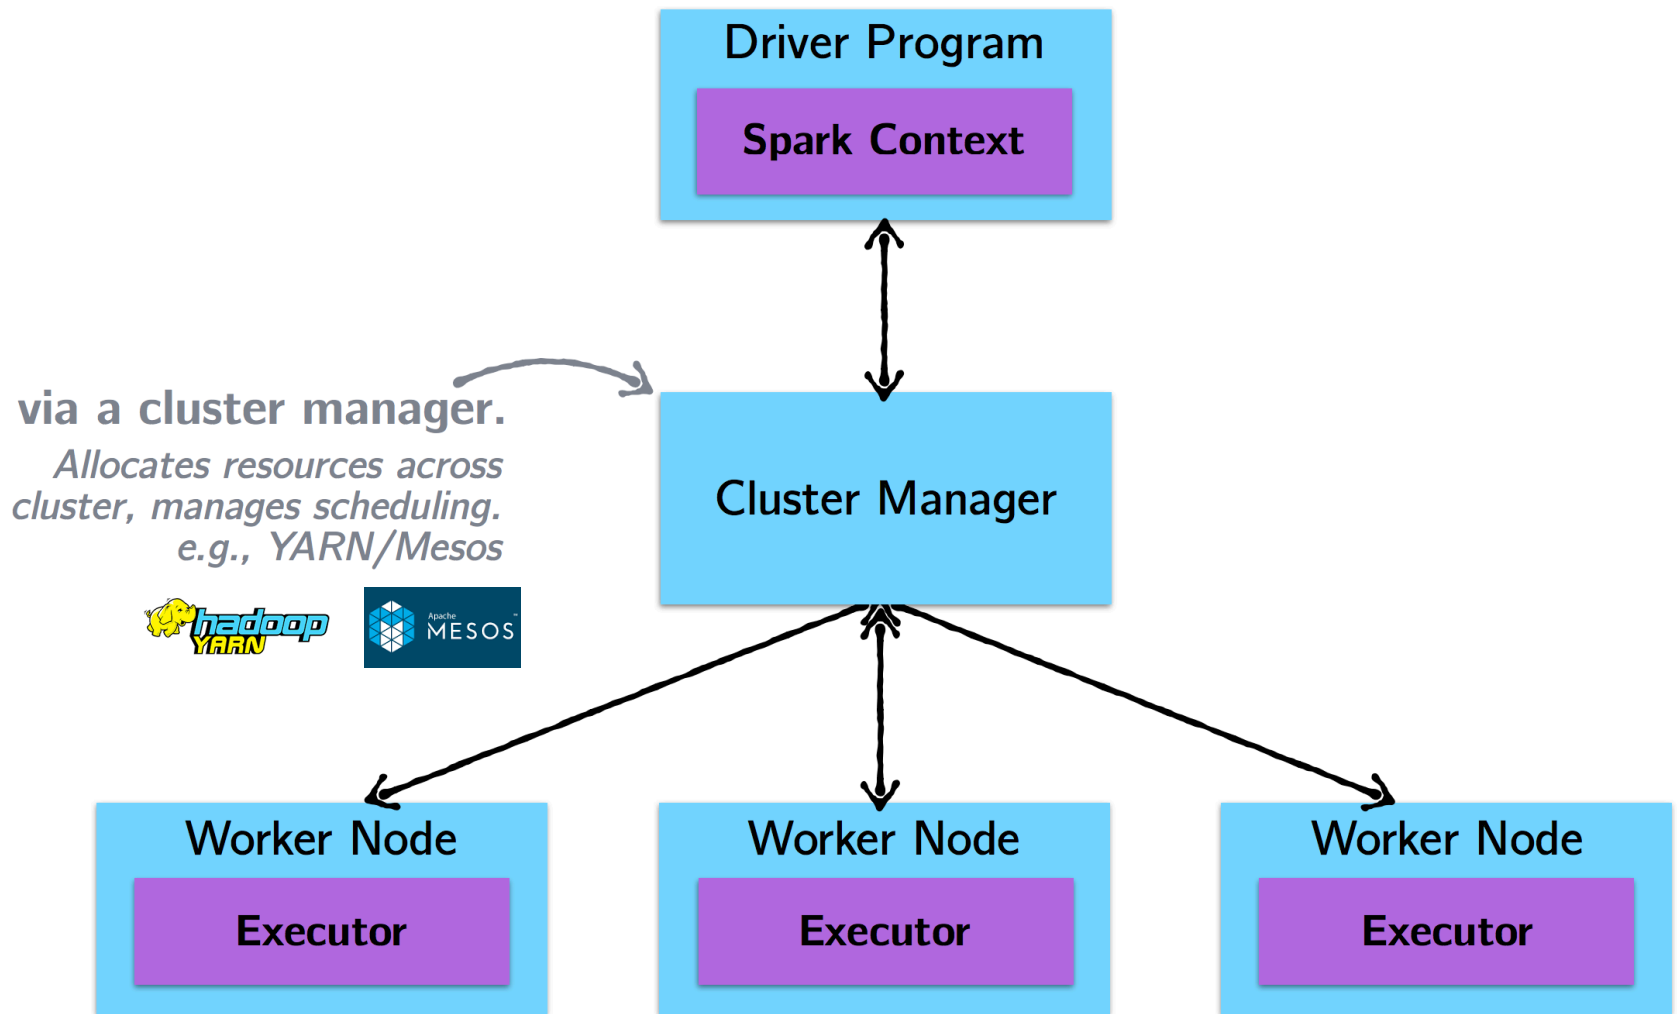
How do they communicate?

These are the nodes actually executing the jobs!



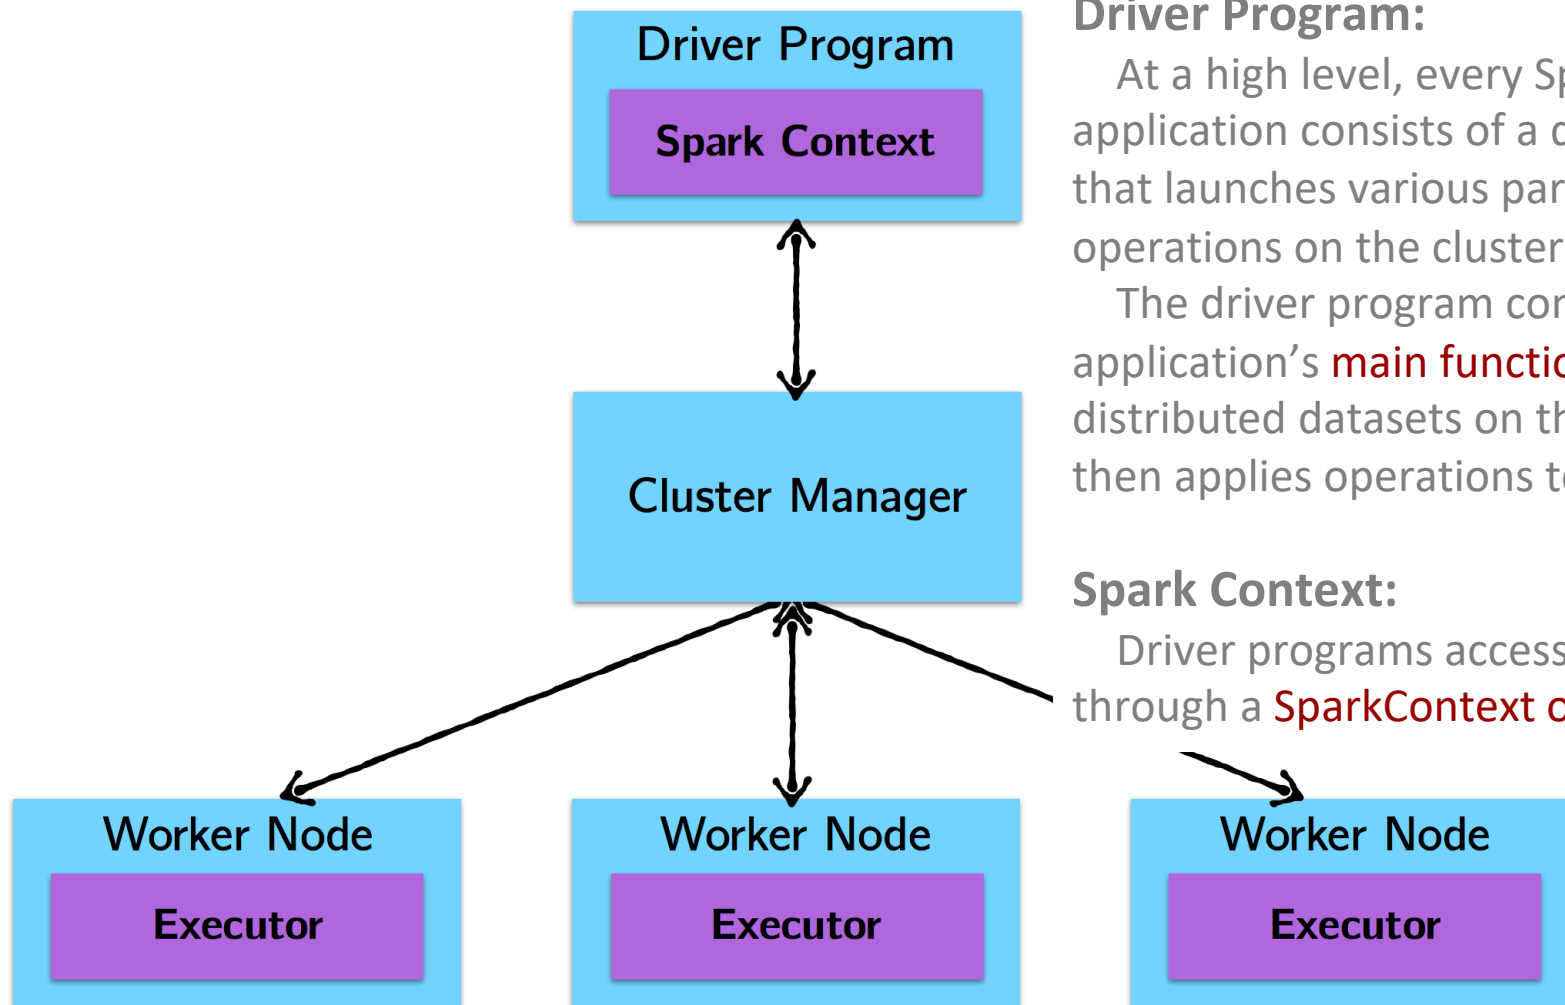


How Spark jobs are Executed





How Spark jobs are Executed



Driver Program:

At a high level, every Spark application consists of a driver program that launches various parallel operations on the cluster.

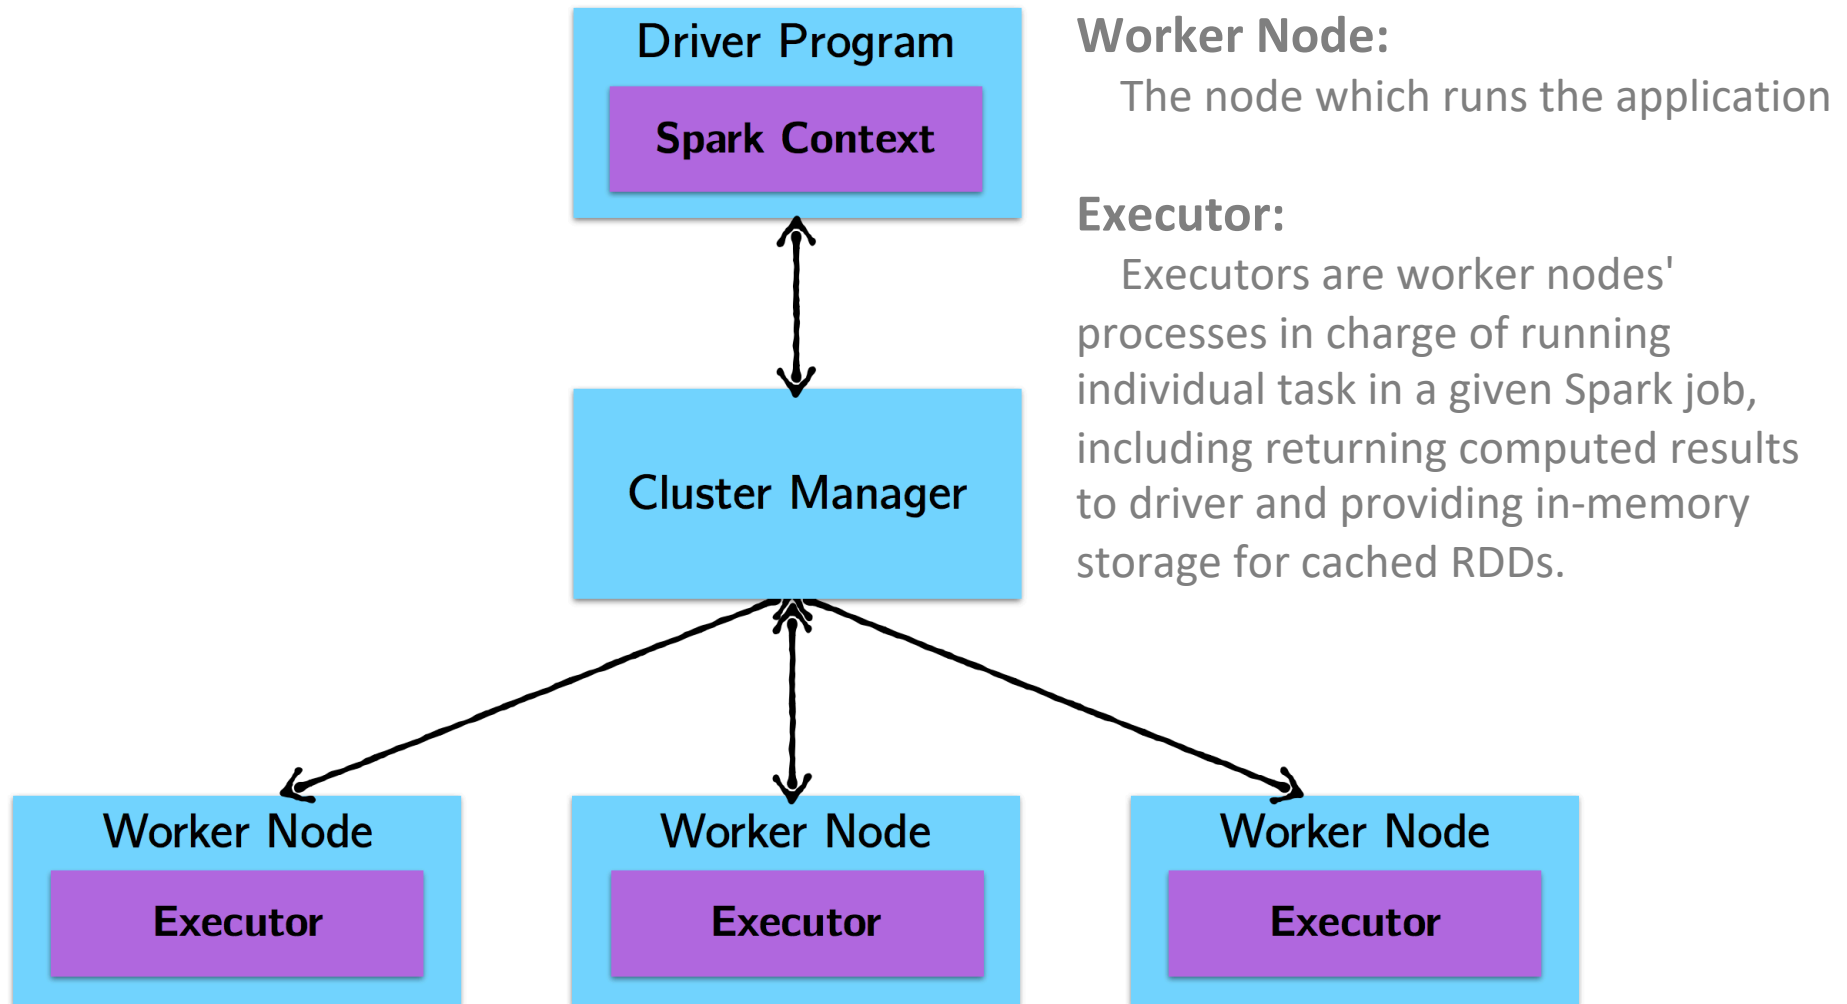
The driver program contains your application's **main function** and defines distributed datasets on the cluster, then applies operations to them.

Spark Context:

Driver programs access Spark through a **SparkContext object**

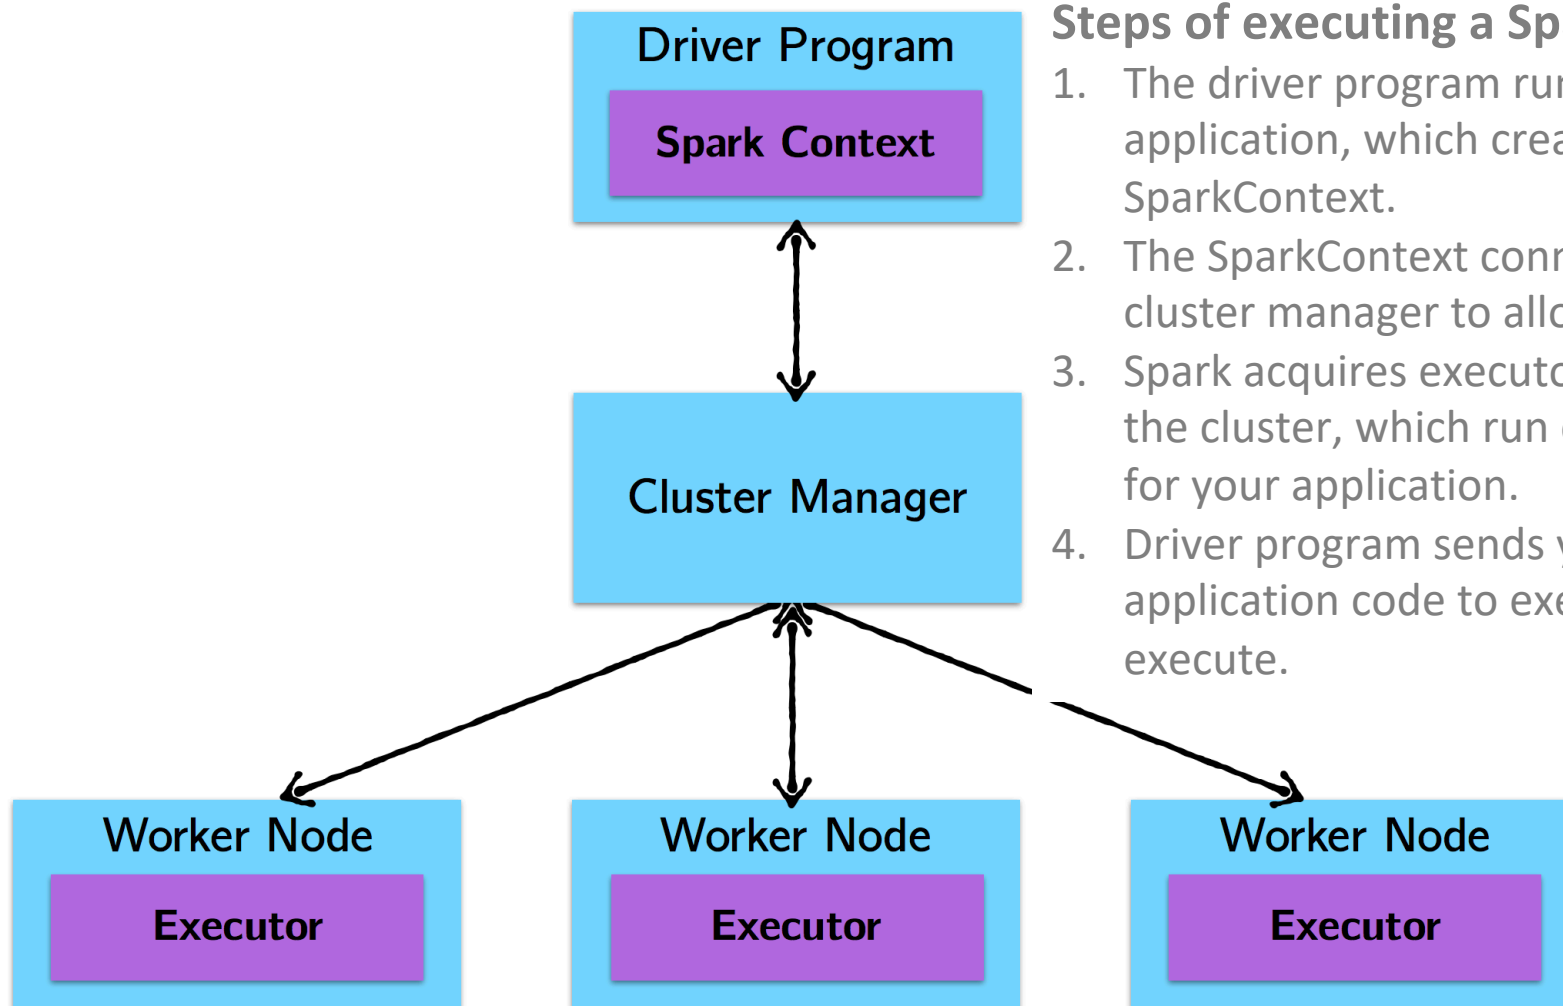


How Spark jobs are Executed





How Spark jobs are Executed



Steps of executing a Spark program:

1. The driver program runs the Spark application, which creates a SparkContext.
2. The SparkContext connects to a cluster manager to allocate resources
3. Spark acquires executors on nodes in the cluster, which run computations for your application.
4. Driver program sends your application code to executors to execute.



Cluster Topology

- A simple example with *println*

```
case class Person(name: String, age: Int)

val people: RDD[Person]
people.foreach(println)
```

What happens?



Cluster Topology

- A simple example with *println*

```
case class Person(name: String, age: Int)

val people: RDD[Person]
people.foreach(println)
```

On the driver: Nothing.

Why? Recall that `foreach` is **an action**, with **return type Unit**. Therefore, it will be eagerly executed on the executors. Thus, any calls to *println* are happening on the worker nodes and are not visible in the driver node.



Cluster Topology

- Another simple example with *take*

```
case class Person(name: String, age: Int)

val people: RDD[Person]
people.foreach(println)

val first10 = people.take(10)
```

Where will *first10* end up?



Cluster Topology

- Another simple example with *take*

```
case class Person(name: String, age: Int)

val people: RDD[Person]
people.foreach(println)

val first10 = people.take(10)
```

Where will *first10* end up? The driver program.

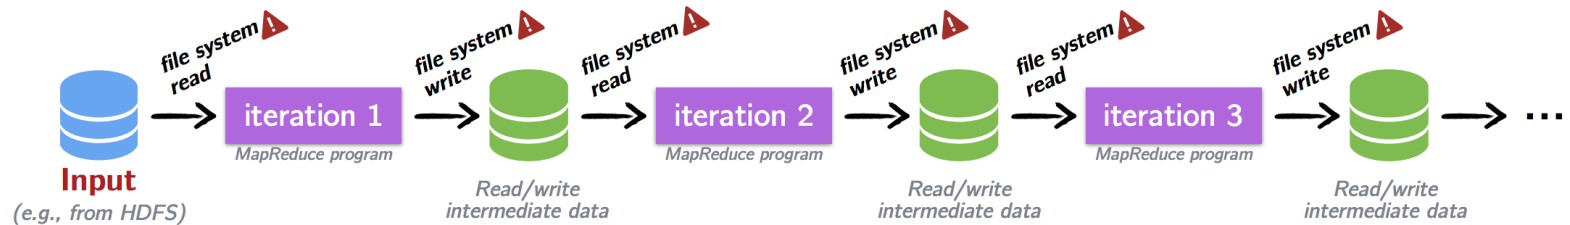
In general, executing an action involves communication between worker nodes and the node running the driver program.



Iteration

- Most data science problems **involve iterations**

Iteration in Hadoop:

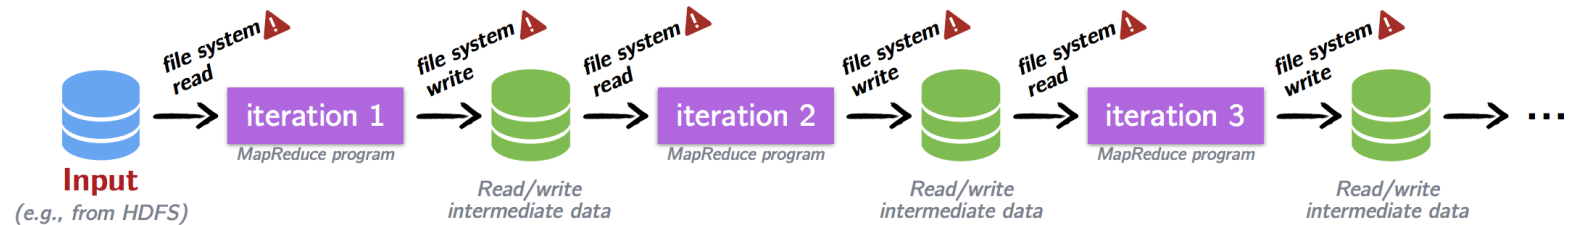





Iteration

- Most data science problems **involve iteration**

Iteration in Hadoop:



Iteration in Spark:






Iteration

Example: Logistic Regression

- Logistic regression is **an iterative algorithm** typically used for classification. Like most classification algorithms, it **updates weights iteratively** base on the training data.





Iteration

Example: Logistic Regression

- Logistic regression is **an iterative algorithm** typically used for classification. Like most classification algorithms, it **updates weights iteratively** base on the training data.

$$w \leftarrow w - \alpha \cdot \sum_{i=1}^n g(w; x_i, y_i)$$



Iteration

Example: Logistic Regression

- Logistic regression sample code:

```
val points = sc.textFile(...).map(parsePoint) // case class Point(x: Double, y: Double)
var w = Vector.zero(d)
for(i <- 1 to numIterations) {
  val gradient = points.map {p =>
    g(p) // Apply the function of logistic regression
  }.reduce(_+_ )
  w -= alpha * gradient
}
```

$$w \leftarrow w - \alpha \cdot \sum_{i=1}^n g(w; x_i, y_i)$$



Iteration

Example: Logistic Regression

- Logistic regression sample code:

```
val points = sc.textFile(...).map(parsePoint) // case class Point(x: Double, y: Double)
var w = Vector.zero(d)
for(i <- 1 to numIterations) {
  val gradient = points.map {p =>
    g(p) // Apply the function of logistic regression
  }.reduce(_+_ )
  w -= alpha * gradient
}
```

$$w \leftarrow w - \alpha \cdot \sum_{i=1}^n g(w; x_i, y_i)$$

What is the weakness for this code?



Iteration

Example: Logistic Regression

- Logistic regression sample code:

```
val points = sc.textFile(...).map(parsePoint) // case class Point(x: Double, y: Double)
var w = Vector.zero(d)
for(i <- 1 to numIterations) {
  val gradient = points.map {p =>
    g(p) // Apply the function of logistic regression
  }.reduce(_+_ )
  w -= alpha * gradient
}
```

Spark starts the execution when the action *reduce* is applied



Iteration

Example: Logistic Regression

- Logistic regression sample code:

```
val points = sc.textFile(...).map(parsePoint) // case class Point(x: Double, y: Double)
var w = Vector.zero(d)
for(i <- 1 to numIterations) {
  val gradient = points.map {p =>
    g(p) // Apply the function of logistic regression
  }.reduce(_+_ )
  w -= alpha * gradient
}
```

***points* is being re-evaluated upon every iteration!**
Unnecessary!



Caching and Persistence

- By default, RDDs are recomputed each time you run an action on them. This can be expensive (time-consuming) if you need to use a dataset more than once.

**Spark allows you to control what is cached in memory
use *persist()* or *cache()***

`cache()` : using the default storage level

`persist()`: can pass the storage level as a parameter,
e.g., “MEMORY_ONLY”, “MEMORY_AND_DISK”



Iteration

Example: Logistic Regression

- Logistic regression sample code:

```
val points = sc.textFile(...).map(parsePoint).persist() // case class Point(x: Double, y: Double)
var w = Vector.zero(d)
for(i <- 1 to numIterations) {
  val gradient = points.map {p =>
    g(p) // Apply the function of logistic regression
  }.reduce(_+_ )
  w -= alpha * gradient
}
```

***points* is evaluated once and is cached in memory.
It can be re-used on each iteration.**

Why Spark is Good for Data Sci



- **The lazy semantics** of RDD transformation operations help improve the performance.
- One of the most common performance bottlenecks for newcomers to Spark arises from unknowingly re-evaluating several transformations when caching could be used.



Other advantages



 **Scala**



python™



- **Easy to use**

- Write applications quickly in Java, Scala, Python

- **Runs Everywhere**

- Spark runs on Hadoop, standalone, or in the cloud
- It can access diverse data sources including HDFS, Cassandra, HBase, and S3



APACHE
HBASE



cassandra

- **Generality**

- Combine SQL, streaming, and complex analytics



amazon
web services™

S3

Word Count

```
word_count.py x text.txt x
1 from pyspark import SparkContext
2 import os
3
4 os.environ['PYSPARK_PYTHON'] = '/usr/local/bin/python3.6'
5 os.environ['PYSPARK_DRIVER_PYTHON'] = '/usr/local/bin/python3.6'
6
7 sc = SparkContext('local[*]', 'wordCount')
8
9 input_file_path = './text.txt'
10 textRDD = sc.textFile(input_file_path)
11
12 counts = textRDD.flatMap(lambda line: line.split(' ')) \
13     .map(lambda word: (word, 1)).reduceByKey(lambda a, b: a+b).collect()
14
15 for each_word in counts:
16     print(each_word)
17
```




If you want to learn more...

- Official documentation
 - <http://spark.apache.org/docs/latest/>
- Online course
 - Coursera: Big Data Analysis with Scala and Spark
- Books
 - *Learning Spark, O' Reilly*
 - *Advanced Analytics with Spark: Patterns for Learning from Data at Scale, O' Reilly*
 - *Machine Learning with Spark, Packt*